

Author: Ricardo Huarte Salazar

About Data Set:

You will find 2 .csv files attached to this task. 1 of the files consist of courier's lifetime dependent features and other consist courier's weekly variant features. Features are renamed for confidentiality purposes and data dictionary will NOT be provided. However, in 2 different .csv files, same courier ID represents same courier.

In [328]:

```
1 pd
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

In [2]:

```
1 %matplotlib inline
```

In [3]:

```
1 lifetime = pd.read_csv(filepath_or_buffer='Courier_lifetime_data.csv')
```

In [4]:

```
1 weekly = pd.read_csv(filepath_or_buffer='Courier_weekly_data.csv')
```

In [5]:

```
1 lifetime.dtypes
```

Out[5]:


```
courier      int64
feature_1    object
feature_2    float64
dtype: object
```

In [6]:

```
1 plt.hist(lifetime['feature_1'])
```

Out[6]:

```
(array([2516.,    0.,    0., 4456.,    0.,    0.,   85.,    0.,    0.,
        467.]),
 array([0. , 0.3, 0.6, 0.9, 1.2, 1.5, 1.8, 2.1, 2.4, 2.7, 3. ]),
 <a list of 10 Patch objects>)
```




In [7]:

```
1 plt.hist(lifetime['feature_2'][~np.isnan(lifetime['feature_2'])])
```

Out[7]:

```
(array([5.915e+03, 6.700e+02, 0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00,
        0.000e+00, 0.000e+00, 0.000e+00, 3.000e+00]),
 array([-61. ,  40.5, 142. , 243.5, 345. , 446.5, 548. , 649.5, 751. ,
        852.5, 954. ]),
 <a list of 10 Patch objects>)
```



In [367]:

```
1 lifetime[(lifetime.feature_2.isnull())].head()
```

Out[367]:

	courier	feature_1	feature_2
2	225	c	NaN
4	242	c	NaN
5	350	a	NaN
6	645	a	NaN
7	1210	a	NaN

In [9]:

```
1 lifetime.describe()
```

Out[9]:

	courier	feature_2
count	7524.000000	6588.000000
mean	518864.440324	26.373862
std	286880.574472	22.703621
min	208.000000	-61.000000
25%	275875.750000	20.000000
50%	529366.500000	25.000000
75%	803120.500000	32.000000
max	964240.000000	954.000000

In [11]:

```
1 lifetime_encoded= pd.get_dummies(lifetime, columns=['feature_1'])
```

In [368]:

```
1 lifetime_encoded.head()
```

Out[368]:


	feature_2	feature_1_a	feature_1_b	feature_1_c	feature_1_d
courier					
208	25.0	1	0	0	0
218	0.0	0	0	1	0
225	NaN	0	0	1	0
231	0.0	0	0	1	0
242	NaN	0	0	1	0

In [13]:

```
1 msno.matrix(lifetime)
```

Out[13]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1dfa1b5fa90>
```




In [18]:

```
1 msno.matrix(weekly)
```

Out[18]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1dfa1c3f3c8>
```



In [19]:

```
1 weekly['courier'].drop_duplicates().count()
```

Out[19]:

759

In [20]:


```
1 weekly[(weekly.week==11) | (weekly.week==10) | (weekly.week==9)][ 'courier'].drop_duplicates().count()
```

Out[20]:

387

In [21]:

```
1 lifetime.hist(bins=50, figsize=(30,20));
```



Task 1: Exploratory Analysis and Data Munging

In this task, you are being expected to clean data, treat missing values, find out related features and finally label the data. Every courier did not work every week. Thus, some of courier-week combinations' data are not provided. First, come up with a way to treat these missing values. Removing missing values are not suggested since provided data set is small and it will affect your predictive model's evaluation metric. Create a report / dashboard and correlation matrix, in addition to results of your univariate and bivariate analysis and explain your findings. Finally, label your data. If a specific courier's week 9, 10 and 11 data is not provided, we label this courier as "1" otherwise "0". After labeling, remove week 8(Yes including 8!), 9, 10 and 11 data to avoid bias in your next task. In addition, distribution of feature_3 is a hint how the data is generated.

In [23]:

1 weekly.describe().T

Out[23]:

	count	mean	std	min	25%	50%
courier	4117.0	366530.934418	128603.611959	3767.000000	280239.000000	406936.000000
week	4117.0	4.910857	3.364852	0.000000	2.000000	5.000000
feature_1	4117.0	-3.702453	17.407331	-138.000000	-12.000000	-2.000000
feature_2	4117.0	44.232208	24.007116	1.000000	26.000000	41.000000
feature_3	4117.0	55.691037	31.666550	1.000000	31.000000	51.000000
feature_4	4117.0	0.068610	0.068999	0.000000	0.018500	0.054100
feature_5	4117.0	0.931390	0.068999	0.000000	0.901200	0.945900
feature_6	4117.0	104.331502	8.473348	92.857100	100.000000	100.465100
feature_7	4117.0	0.059339	0.064646	0.000000	0.000000	0.043500
feature_8	4117.0	3975.807328	1237.055134	1136.750000	2750.977800	4099.425000
feature_9	4117.0	0.767527	0.136458	0.000000	0.693700	0.785700
feature_10	4117.0	9.619359	1.827863	2.575000	8.424751	9.497961
feature_11	4117.0	20.266942	12.460020	0.000000	11.000000	19.000000
feature_12	4117.0	20.000994	3.205479	5.416667	18.168824	19.648810
feature_13	4117.0	5.211435	0.961980	3.270000	4.570099	5.072500
feature_14	4117.0	0.782381	0.164578	0.000000	0.739100	0.822200
feature_15	4117.0	68.655642	18.828885	2.957809	57.839947	71.653595
feature_16	4117.0	2.255526	1.542969	1.000000	1.000000	2.000000
feature_17	4117.0	12.789410	11.691080	1.000000	5.000000	10.000000

In [369]:

1 weekly.head()

Out[369]:

	courier	week	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6	feature_7	feature_8
0	3767	2	6	34	38	0.0789	0.9211	140.4737	0.1316	216
1	3767	4	-1	42	37	0.0000	1.0000	135.5946	0.0811	209
2	3767	5	24	41	43	0.0233	0.9767	131.0930	0.0233	204
3	3767	6	-22	65	66	0.0606	0.9394	120.1515	0.0000	211
4	6282	2	9	33	27	0.0741	0.9259	100.0000	0.0370	407

In []:

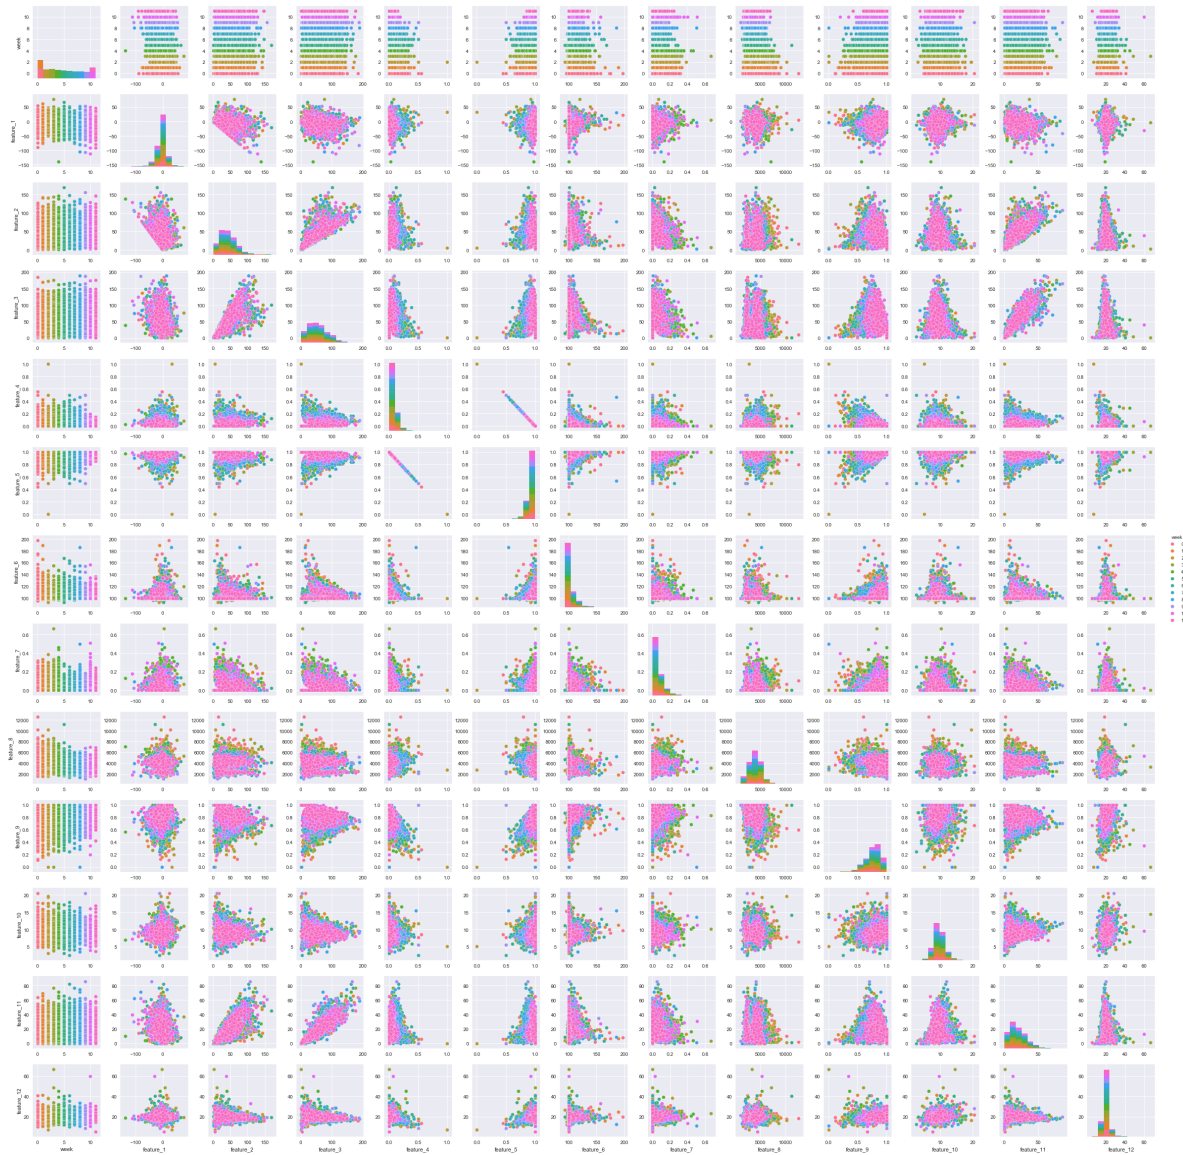
1

In [25]:

```
1 sns.pairplot(weekly[['week', 'feature_1', 'feature_2', 'feature_3', 'feature_4', 'feature_5
```

Out[25]:

<seaborn.axisgrid.PairGrid at 0x1dfa4389a58>



In [26]:


```
1 analysis=weekly[['week', 'feature_1', 'feature_2', 'feature_3', 'feature_4', 'feature_5', 'f
2 corr_mt=analysis.corr()
```

In [27]:

```
1 corr_mt
```

Out[27]:

	week	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6	feature_7
week	1.000000	-0.084761	0.098734	0.095960	-0.206059	0.206057	-0.059651	-0.253702
feature_1	-0.084761	1.000000	-0.212915	-0.132828	-0.029416	0.029420	0.063833	0.056426
feature_2	0.098734	-0.212915	1.000000	0.788146	0.188848	-0.188848	0.039640	-0.003138
feature_3	0.095960	-0.132828	0.788146	1.000000	0.162444	-0.162446	0.062238	0.027473
feature_4	-0.206059	-0.029416	0.188848	0.162444	1.000000	-1.000000	-0.073383	0.038013
feature_5	0.206057	0.029420	-0.188848	-0.162446	-1.000000	1.000000	0.073384	-0.038012
feature_6	-0.059651	0.063833	0.039640	0.062238	-0.073383	0.073384	1.000000	0.162079
feature_7	-0.253702	0.056426	-0.003138	0.027473	0.038013	-0.038012	0.162079	1.000000
feature_8	-0.192313	-0.008650	0.133437	0.068195	0.227298	-0.227299	-0.144915	0.117473
feature_9	0.302481	0.057826	-0.152309	-0.103423	-0.636237	0.636239	0.148237	0.001861
feature_10	-0.165839	0.062050	-0.146601	-0.173151	-0.175879	0.175880	-0.016608	-0.002511
feature_11	0.042760	-0.100593	0.749931	0.863150	0.073725	-0.073725	0.061017	0.019581
feature_12	-0.028613	-0.004104	-0.183353	-0.179758	-0.285753	0.285752	0.019124	-0.013851
feature_13	-0.189233	-0.013375	0.012738	-0.060513	0.115765	-0.115766	0.242932	-0.012011
feature_14	0.223040	0.062439	0.010894	0.063962	-0.454770	0.454773	0.129238	0.096681
feature_15	-0.205423	0.124982	-0.162611	0.316081	0.047795	-0.047795	0.030601	0.187241
feature_16	0.055816	-0.091912	0.430875	0.505898	0.032682	-0.032686	0.031249	0.097791
feature_17	-0.107284	-0.139890	0.396970	0.421079	0.171434	-0.171435	0.020535	0.046561




In [28]:

```
1 sns.heatmap(corr_mt, vmax=1., square=False, cmap="RdYlBu")
```

Out[28]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1dfa4bfa208>
```



In []:

```
1
```

In [29]:

```
1 weekly.head()
```

Out[29]:

	courier	week	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6	feature_7	feature_8
0	3767	2	6	34	38	0.0789	0.9211	140.4737	0.1316	216
1	3767	4	-1	42	37	0.0000	1.0000	135.5946	0.0811	209
2	3767	5	24	41	43	0.0233	0.9767	131.0930	0.0233	204
3	3767	6	-22	65	66	0.0606	0.9394	120.1515	0.0000	211
4	6282	2	9	33	27	0.0741	0.9259	100.0000	0.0370	401

In []:

```
1
```

In [30]:

```
1 week=weekly.copy()
```

In [31]:

```
1 def week_label(row):
2     courier_set=weekly[(weekly.courier==row['courier']) & ((weekly.week==9) | (weekly.week==10))]
3     if courier_set['courier'].count() == 0:
4         label=1
5     else:
6         label=0
7     return label
```

In [32]:

```
1 week['label']=week.apply(week_label, axis=1)
```

In [370]:

```
1 week.head()
```

Out[370]:

	courier	week	feature_1	feature_2	feature_3	feature_4	feature_6	feature_7	feature_8	feature_9
0	3767	2	6	34	38	0.0789	140.4737	0.1316	2162.4737	0.0000
1	3767	4	-1	42	37	0.0000	135.5946	0.0811	2097.4054	0.0000
2	3767	5	24	41	43	0.0233	131.0930	0.0233	2043.8837	0.0000
3	3767	6	-22	65	66	0.0606	120.1515	0.0000	2124.2727	0.0000
4	6282	2	9	33	27	0.0741	100.0000	0.0370	4075.7407	0.0000

In [34]:

```
1 corr_mt=week.corr()
2 corr_mt
```

Out[34]:


	courier	week	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6
courier	1.000000	-0.024238	0.016642	-0.276750	-0.093481	-0.119213	0.119214	-0.074533
week	-0.024238	1.000000	-0.084761	0.098734	0.095960	-0.206059	0.206057	-0.059651
feature_1	0.016642	-0.084761	1.000000	-0.212915	-0.132828	-0.029416	0.029420	0.063833
feature_2	-0.276750	0.098734	-0.212915	1.000000	0.788146	0.188848	-0.188848	0.039640
feature_3	-0.093481	0.095960	-0.132828	0.788146	1.000000	0.162444	-0.162446	0.062238
feature_4	-0.119213	-0.206059	-0.029416	0.188848	0.162444	1.000000	-1.000000	-0.073383
feature_5	0.119214	0.206057	0.029420	-0.188848	-0.162446	-1.000000	1.000000	0.073383
feature_6	-0.074533	-0.059651	0.063833	0.039640	0.062238	-0.073383	0.073384	1.000000
feature_7	-0.016657	-0.253702	0.056426	-0.003138	0.027473	0.038013	-0.038012	0.162071
feature_8	-0.115310	-0.192313	-0.008650	0.133437	0.068195	0.227298	-0.227299	-0.144911
feature_9	0.015373	0.302481	0.057826	-0.152309	-0.103423	-0.636237	0.636239	0.148237
feature_10	0.086274	-0.165839	0.062050	-0.146601	-0.173151	-0.175879	0.175880	-0.016601
feature_11	-0.129342	0.042760	-0.100593	0.749931	0.863150	0.073725	-0.073725	0.061017
feature_12	0.058263	-0.028613	-0.004104	-0.183353	-0.179758	-0.285753	0.285752	0.019121
feature_13	-0.046374	-0.189233	-0.013375	0.012738	-0.060513	0.115765	-0.115766	0.242931
feature_14	-0.035208	0.223040	0.062439	0.010894	0.063962	-0.454770	0.454773	0.129231
feature_15	0.161679	-0.205423	0.124982	-0.162611	0.316081	0.047795	-0.047795	0.030601
feature_16	0.010081	0.055816	-0.091912	0.430875	0.505898	0.032682	-0.032686	0.031241
feature_17	-0.002531	-0.107284	-0.139890	0.396970	0.421079	0.171434	-0.171435	0.020531
label	0.051439	-0.307316	-0.101587	-0.294234	-0.349718	0.034596	-0.034597	-0.031041

In [35]:

```
1 sns.heatmap(corr_mt, vmax=1., square=False, cmap="RdYlBu")
```

Out[35]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1dfa9c53668>
```



In [284]:

```
1 week.drop('feature_5', axis=1, inplace=True)
2 week.drop('feature_11', axis=1, inplace=True)
```

if we look at the correlation matrix we can see that feature 4 and 5 are highly correlated, hence we will get rid of the column 5, we also see that 3 and eleven are have also a high correlation coefficient

In [285]:

```
1 week_2=week[(week.week<8)]
```

In [287]:

```
1 transposed= week_2.pivot(index='courier', columns='week')
```

In [288]:

```
1 transposed.columns.get_level_values(0)
```

Out[288]:

```
Index(['feature_1', 'feature_1', 'feature_1', 'feature_1', 'feature_1',
      'feature_1', 'feature_1', 'feature_1', 'feature_2', 'feature_2',
      ...,
      'feature_17', 'feature_17', 'label', 'label', 'label', 'label', 'label',
      'label', 'label', 'label'],
      dtype='object', length=128)
```

In [289]:

```
1 transposed.columns = [transposed + '_' + i for transposed, i in zip(transposed.columns
```

In []:

```
1 lifetime_encoded.set_index('courier', inplace=True)
```

In [290]:

```
1 transposed=transposed.merge(lifetime_encoded, how='inner', left_index=True, right_index=True)
```

In [291]:

```
1 transposed.head()
```

Out[291]:

	feature_1_0	feature_1_1	feature_1_2	feature_1_3	feature_1_4	feature_1_5	feature_1_6
courier							
3767	NaN	NaN	6.0	NaN	-1.0	24.0	-22.0
6282	NaN	NaN	9.0	-20.0	9.0	21.0	-12.0
10622	5.0	-12.0	NaN	NaN	NaN	NaN	NaN
13096	NaN	NaN	NaN	NaN	NaN	-10.0	10.0
14261	4.0	-16.0	2.0	3.0	7.0	-1.0	-1.0

5 rows × 133 columns

In []:

```
1
```

In [371]:

```
1 transposed_mt=transposed.corr()
2 transposed_mt.head()
```

Out[371]:

	feature_1_0	feature_1_1	feature_1_2	feature_1_3	feature_1_4	feature_1_5	feature_1_6
feature_1_0	1.000000	-0.251950	-0.029288	-0.082137	-0.117810	0.038922	0.082137
feature_1_1	-0.251950	1.000000	-0.374919	0.128965	0.211154	0.001474	-0.029288
feature_1_2	-0.029288	-0.374919	1.000000	-0.267218	0.007909	0.114968	-0.007909
feature_1_3	-0.082137	0.128965	-0.267218	1.000000	-0.289214	0.115630	-0.051474
feature_1_4	-0.117810	0.211154	0.007909	-0.289214	1.000000	-0.347923	0.062137


5 rows × 126 columns

In [293]:

```
1 sns.heatmap(transposed_mt, vmax=1., square=False, cmap="RdYlBu")
```

Out[293]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1dfba30d2b0>
```



In [294]:

```
1 def get_label(row):
2     label=0
3     for i in range(8):
4         column='label_'+ str(i)
5         if row[column]==1:
6             label=1
7     return label
8
```

In [295]:

```
1 transposed['y']=transposed.apply(get_label,axis=1)
```

In [296]:

```
1 columns_d1=['label_0','label_1','label_2','label_3','label_4','label_5','label_6','lab
```

In [297]:

```
1 transposed.drop(['label_0','label_1','label_2','label_3','label_4','label_5','label_6'
```

In [372]:

```
1 transposed.head()
```

Out[372]:

	feature_1_0	feature_1_1	feature_1_2	feature_1_3	feature_1_4	feature_1_5	feature_1_6
courier							
3767	NaN	NaN	6.0	NaN	-1.0	24.0	-22.0
6282	NaN	NaN	9.0	-20.0	9.0	21.0	-12.0
10622	5.0	-12.0	NaN	NaN	NaN	NaN	NaN
13096	NaN	NaN	NaN	NaN	NaN	-10.0	10.0
14261	4.0	-16.0	2.0	3.0	7.0	-1.0	-1.0

5 rows × 126 columns


Imputing Missing values using KNN and FancyImpute

In [299]:

```
1 msno.matrix(transposed)
```

Out[299]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1dfba39ab38>
```



I chose to use KNN imputation because I feel like it gives us a better idea of the users than other imputation methods like means and modes,

I also tried imputing zero (0) to the features where we had no information but that approach gave me less predicting power than KNN Imputation

In [300]:

```
1 transposed_filled = pd.DataFrame(KNN(k=2).fit_transform(transposed), columns=transposed)
```


```
Imputing row 1/729 with 60 missing, elapsed time: 0.896
Imputing row 101/729 with 105 missing, elapsed time: 1.022
Imputing row 201/729 with 15 missing, elapsed time: 1.138
Imputing row 301/729 with 60 missing, elapsed time: 1.366
Imputing row 401/729 with 105 missing, elapsed time: 1.533
Imputing row 501/729 with 0 missing, elapsed time: 1.700
Imputing row 601/729 with 90 missing, elapsed time: 1.838
Imputing row 701/729 with 105 missing, elapsed time: 2.015
```

In [301]:

```
1 msno.matrix(transposed_filled)
```

Out[301]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1dfba46ee10>
```



Predicting Classes

Labels are almost balance at 50/50

In [302]:

```
1 golden_x= transposed_filled.drop(['y'], axis=1)
2 golden_y= transposed_filled['y']
```

In [303]:

```
1 x_train, x_test, y_train, y_test = train_test_split(golden_x, golden_y , test_size=0.25)
```


In [304]:

```
1 golden_y.value_counts()
```

Out[304]:

```
0.0    365
```

```
1.0    364
```

```
Name: y, dtype: int64
```

Using Logistic Regression w Grid Search

In [362]:

```
1 clf = LogisticRegression()
```

In [334]:

```

1 logisticRegr = LogisticRegression()
2 logisticRegr.fit(X=x_train, y=y_train)
3 test_y_pred = logisticRegr.predict(x_test)
4 cf_mt = confusion_matrix(y_test, test_y_pred)
5 print('Intercept: ' + str(logisticRegr.intercept_))
6 print('Regression: ' + str(logisticRegr.coef_))
7 print('Accuracy of logistic regression classifier on test set: {:.2f}'.format(logistic
8 print(classification_report(y_test, test_y_pred))

```

Intercept: [0.08250367]

Regression: [[-2.33585280e-02 2.23843424e-02 2.90318108e-03 -4.97461652e-0

2

```

-4.55416559e-02 -4.00217777e-02 -5.01389168e-02 -6.51825045e-02
 2.12225610e-02 -4.08226525e-03 -4.50195507e-03 -5.02704635e-02
 4.06239407e-02 -1.04078398e-02  2.42477353e-02 -4.49304715e-02
-2.01439034e-02 -3.31138091e-02  2.85915143e-02  2.64477053e-02
-3.87853683e-02  6.32579392e-03 -3.50290578e-02  1.54860836e-02
 2.72099330e-02  3.31997524e-01  1.55754312e-01  2.13096596e-01
-8.74812485e-02  8.35084541e-02  3.55258397e-01  2.95026245e-01
-4.64646054e-02  3.30660486e-02  1.19925375e-01 -6.86550009e-03
 3.45330504e-02 -2.26436556e-02 -3.50423441e-02 -1.35494007e-01
-1.54684973e-01 -7.45526199e-02 -4.89261942e-02 -1.75692309e-01
-2.45132022e-01 -2.19442584e-01 -8.96292455e-02  2.85270621e-03
-3.10961885e-04  1.02242770e-03  7.04419029e-04 -5.15864045e-04
-5.24781319e-04 -2.21963170e-04  1.79660613e-04  6.22287164e-04
 9.93032797e-02  1.88471722e-01 -1.18169809e-01 -5.26891529e-01
-1.12227822e-01 -2.52203622e-01 -2.40071696e-01 -4.41510958e-01
 2.06468724e-01 -1.62731958e-01  1.37459154e-01  9.23105767e-02
 1.27127652e-01  1.83976451e-01  2.35317682e-01 -2.67812355e-01
-1.49208490e-03  1.60283144e-01  1.52228973e-01 -5.17015997e-02
 4.79080321e-02  3.96172661e-03 -1.35737920e-01 -2.09202164e-02
 4.99636533e-01 -6.79362118e-01 -1.60576250e+00  2.22487018e-01
 5.69455213e-01  4.24346880e-01  1.05381649e+00  2.60465940e-01
-4.45071530e-02  1.42882516e-01  9.96511914e-02 -3.25139684e-01
 5.22873426e-01 -4.20723472e-02 -4.09887397e-01 -3.20157612e-01
 1.44207259e-02  4.52463544e-05 -2.03665615e-02  1.42798672e-03
 6.03142835e-03 -8.73470413e-03  1.16894701e-02 -3.53672933e-02
-9.49074407e-03  1.64685921e-01 -9.31360653e-01  3.56551548e-01
-4.09565816e-01 -7.27424496e-01 -5.29651816e-01 -2.31923629e-01
-4.08153598e-02 -1.94702323e-02  5.71749753e-02 -2.44126983e-02
 7.09380087e-02 -2.32544502e-02 -4.96074572e-02 -8.47006029e-02
 9.82610436e-03  8.81166869e-02 -1.49438433e-01  2.81994134e-01
-1.38168714e-01]]

```

Accuracy of logistic regression classifier on test set: 0.79


	precision	recall	f1-score	support
0.0	0.79	0.78	0.78	89
1.0	0.79	0.81	0.80	94
avg / total	0.79	0.79	0.79	183

In [335]:

```
1 confusion_matrix_df = pd.DataFrame(cf_mt, ('False', 'True'), ('False', 'True'))
2 heatmap = sns.heatmap(confusion_matrix_df, annot=True, annot_kws={"size": 20}, fmt="d")
3 heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right', fo
4 heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45, ha='right',
5 plt.ylabel('True label', fontsize = 14)
6 plt.xlabel('Predicted label', fontsize = 14)
```

Out[335]:

Text(0.5,16,'Predicted label')



In [336]:

```

1  clf.fit(x_train,y_train)
2  print(clf.score(x_test, y_test))
3  seed = 7
4  k_fold = KFold(n_splits=10, random_state=seed)
5  scoring = 'accuracy'
6  results = results=cross_val_score(clf, x_test, y_test, cv=k_fold, n_jobs=1, scoring=scoring)
7  print("Accuracy: %.3f (%.3f)" % (results.mean(), results.std()))
8  scoring = 'neg_log_loss'
9  results = results=cross_val_score(clf, x_test, y_test, cv=k_fold, n_jobs=1, scoring=scoring)
10 print("Logloss: %.3f (%.3f)" % (results.mean(), results.std()))
11 scoring = 'roc_auc'
12 results = results=cross_val_score(clf, x_test, y_test, cv=k_fold, n_jobs=1, scoring=scoring)
13 print("AUC: %.3f (%.3f)" % (results.mean(), results.std()))
14
15 clf.fit(x_train,y_train)
16 predicted=clf.predict(x_test)
17 matrix = confusion_matrix(y_test, predicted)
18 print(matrix)
19
20
21 confusion_matrix_df = pd.DataFrame(matrix, ('False', 'True'), ('False', 'True'))
22 heatmap = sns.heatmap(confusion_matrix_df, annot=True, annot_kws={"size": 20}, fmt="d",
23                        yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right', fontweight='bold'),
24                        xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45, ha='right', fontweight='bold'),
25                        plt.ylabel('True label', fontsize = 14)
26                        plt.xlabel('Predicted label', fontsize = 14)
27
28

```

0.7923497267759563

Accuracy: 0.672 (0.135)

Logloss: -4.142 (1.447)


AUC: 0.729 (0.162)

[[69 20]

[18 76]]


Out[336]:

Text(0.5,16,'Predicted label')



In [337]:

```
1 probs = clf.predict_proba(x_test)
2 preds = probs[:,1]
3 fpr, tpr, threshold = roc_curve(y_test, preds)
4 roc_auc = auc(fpr, tpr)
5
6 precis
7
8 plt.title('Receiver Operating Characteristic')
9 plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
10 plt.legend(loc = 'lower right')
11 plt.plot([0, 1], [0, 1], 'r--')
12 plt.xlim([0, 1])
13 plt.ylim([0, 1])
14 plt.ylabel('True Positive Rate')
15 plt.xlabel('False Positive Rate')
16 plt.show()
```



I Used Randomize search with Logistic regression since with fewer iterations and less time is more likely to find the optimal parameters

In [363]:

```

1 penalty = ['l2']
2 #C = np.logspace(0, 4, 10)
3 C = uniform(loc=0, scale=4)
4 hyperparameters = dict(C=C, penalty=penalty)
5
6 #clf_CV = GridSearchCV(clf, hyperparameters, cv=5, verbose=0)
7 clf_CV = RandomizedSearchCV(clf, hyperparameters, random_state=1, n_iter=10, cv=5, ver
8 best_model = clf_CV.fit(x_train, y_train)
9 clf_CV.fit(x_train,y_train)
10 print(clf_CV.score(x_test, y_test))
11 seed = 7
12 k_fold = KFold(n_splits=10, random_state=seed)
13 scoring = 'accuracy'
14 results = results=cross_val_score(clf_CV, x_test, y_test, cv=k_fold, n_jobs=1, scoring:
15 print("Accuracy: %.3f (%.3f)" % (results.mean(), results.std()))
16 scoring = 'neg_log_loss'
17 results = results=cross_val_score(clf_CV, x_test, y_test, cv=k_fold, n_jobs=1, scoring:
18 print("Logloss: %.3f (%.3f)" % (results.mean(), results.std()))
19 scoring = 'roc_auc'
20 results = results=cross_val_score(clf_CV, x_test, y_test, cv=k_fold, n_jobs=1, scoring:
21 print("AUC: %.3f (%.3f)" % (results.mean(), results.std()))
22
23 predicted=best_model.predict(x_test)
24 matrix = confusion_matrix(y_test, predicted)
25
26
27 confusion_matrix_df = pd.DataFrame(matrix, ('False', 'True'), ('False', 'True'))
28 heatmap = sns.heatmap(confusion_matrix_df, annot=True, annot_kws={"size": 20}, fmt="d"
29 heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right', f
30 heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45, ha='right',
31 plt.ylabel('True label', fontsize = 14)
32 plt.xlabel('Predicted label', fontsize = 14)
33

```

0.8032786885245902


Accuracy: 0.744 (0.082)

Logloss: -0.592 (0.160)

AUC: 0.811 (0.098)

Out[363]:

Text(0.5,16,'Predicted label')




In [366]:

```

1 probs = best_model.predict_proba(x_test)
2 preds = probs[:,1]
3 fpr, tpr, threshold = roc_curve(y_test, preds)
4 roc_auc = auc(fpr, tpr)
5
6 print(accuracy_score(y_test, predicted))
7
8 plt.title('Receiver Operating Characteristic')
9 plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
10 plt.legend(loc = 'lower right')
11 plt.plot([0, 1], [0, 1], 'r--')
12 plt.xlim([0, 1])
13 plt.ylim([0, 1])
14 plt.ylabel('True Positive Rate')
15 plt.xlabel('False Positive Rate')
16 plt.show()

```

0.8032786885245902



I've chosen to use the ROC Curve since we have balanced classes, in this curve we can see how "powerful" our model is, and how false positives are related to true positives, the higher true positive values with a minimum rate of false positive values the better, and we can see in the case of Random Forest, the curve has a good shape and the AUC is higher

In [138]:

1

In []:

1

Using Random Forest Classifier

To complement the prediction part, we try as well the RandomForest algorithm which after parameter tuning yields better results than logistic regression.

In [347]:

```

1 clf = RandomForestClassifier(n_jobs=-1)
2
3 param_grid = {
4     'min_samples_split': [3, 5, 10],
5     'n_estimators' : [100, 300],
6     'max_depth': [3, 5, 15, 25],
7     'max_features': [3, 5, 10, 20]
8 }
9
10 scorers = {
11     'precision_score': make_scorer(precision_score),
12     'recall_score': make_scorer(recall_score),
13     'accuracy_score': make_scorer(accuracy_score)
14 }

```

In [348]:

```

1 def grid_search_wrapper(refit_score='precision_score'):
2     """
3     fits a GridSearchCV classifier using refit_score for optimization
4     prints classifier performance metrics
5     """
6     skf = StratifiedKFold(n_splits=10)
7     grid_search = GridSearchCV(clf, param_grid, scoring=scorers, refit=refit_score,
8                               cv=skf, return_train_score=True, n_jobs=-1)
9     grid_search.fit(x_train.values, y_train.values)
10
11
12     y_pred = grid_search.predict(x_test.values)
13
14     print('Best params for {}'.format(refit_score))
15     print(grid_search.best_params_)
16
17     # confusion matrix on the test data.
18     print('\nConfusion matrix of Random Forest optimized for {} on the test data:'.format(refit_score))
19     print(pd.DataFrame(confusion_matrix(y_test, y_pred),
20                       columns=['pred_neg', 'pred_pos'], index=['neg', 'pos']))
21     return grid_search

```

In [349]:

```
1 grid_search_clf = grid_search_wrapper(refit_score='precision_score')
```

Best params for precision_score

```
{'max_depth': 25, 'max_features': 3, 'min_samples_split': 10, 'n_estimators': 300}
```

Confusion matrix of Random Forest optimized for precision_score on the test data:

	pred_neg	pred_pos
neg	76	13
pos	11	83

In []:

```
1
```


In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In [350]:

```
1 randomForest = RandomForestClassifier( max_depth=25, max_features=3, min_samples_split=
2 randomForest.fit(x_train, y_train)
3 print('Accuracy of random forest classifier on test set: {:.2f}'.format(randomForest.s
```

Accuracy of random forest classifier on test set: 0.86

In [351]:

```
1 test_y_pred = randomForest.predict(x_test)
2 cf_mt = confusion_matrix(y_test, test_y_pred)
3 cf_mt
```

Out[351]:

```
array([[75, 14],
       [11, 83]], dtype=int64)
```

In [357]:

```

1
2 print(randomForest.score(x_test, y_test))
3 seed = 7
4 k_fold = KFold(n_splits=10, random_state=seed)
5 scoring = 'accuracy'
6 results = results=cross_val_score(randomForest, x_test, y_test, cv=k_fold, n_jobs=1, s
7 print("Accuracy: %.3f (%.3f)" % (results.mean(), results.std()))
8 scoring = 'neg_log_loss'
9 results = results=cross_val_score(randomForest, x_test, y_test, cv=k_fold, n_jobs=1, s
10 print("Logloss: %.3f (%.3f)" % (results.mean(), results.std()))
11 scoring = 'roc_auc'
12 results = results=cross_val_score(randomForest, x_test, y_test, cv=k_fold, n_jobs=1, s
13 print("AUC: %.3f (%.3f)" % (results.mean(), results.std()))
14
15 randomForest.fit(x_train,y_train)
16 predicted=randomForest.predict(x_test)
17 matrix = confusion_matrix(y_test, predicted)
18 print(matrix)
19
20
21 confusion_matrix_df = pd.DataFrame(matrix, ('False', 'True'), ('False', 'True'))
22 heatmap = sns.heatmap(confusion_matrix_df, annot=True, annot_kws={"size": 20}, fmt="d"
23 heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right', f
24 heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45, ha='right',
25 plt.ylabel('True label', fontsize = 14)
26 plt.xlabel('Predicted label', fontsize = 14)
27

```

0.8688524590163934

Accuracy: 0.820 (0.063)

Logloss: -0.501 (0.041)


AUC: 0.907 (0.053)

[[74 15]

[13 81]]

Out[357]:

Text(0.5,16,'Predicted label')



In [358]:

```
1 probs = randomForest.predict_proba(x_test)
2 preds = probs[:,1]
3 fpr, tpr, threshold = roc_curve(y_test, preds)
4 roc_auc = auc(fpr, tpr)
5
6
7 plt.title('Receiver Operating Characteristic')
8 plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
9 plt.legend(loc = 'lower right')
10 plt.plot([0, 1], [0, 1], 'r--')
11 plt.xlim([0, 1])
12 plt.ylim([0, 1])
13 plt.ylabel('True Positive Rate')
14 plt.xlabel('False Positive Rate')
15 plt.show()
```



In []:

1